

QUORA INSINCERE CLASSIFY

Nguyễn Thành Long & Nguyễn Bá Nghĩa

Introduction & Overview of the problem:

An existential problem for any major website today is how to handle toxic and divisive content. Quora is a platform to gain & share knowledge where you can ask any question and get answers from different people with unique insights. At the same time, it's important to handle the toxic contents to make the users safe to share their knowledge.

Dataset the Task:

The task: Build the model to predict whether a question asked is sincere or insincere

Overview about dataset:

Link to download dataset:

<https://www.kaggle.com/c/quora-insincere-questions-classification/data>


About this dataset:

- Embeddings.zip: Some pretrain embedding set (Glove, Word2Vec, FastText, Paragram)
- Train.csv: Training set with 1306122 samples is question that labeled to 0 (sincere) or 1 (insincere).
- Test.csv: Testing set includes 56370 samples is question without label

Because we can't evaluate on test set, we decide to divide training set up to three subsets: train set, validation set and test set with ratio 8:1:1.

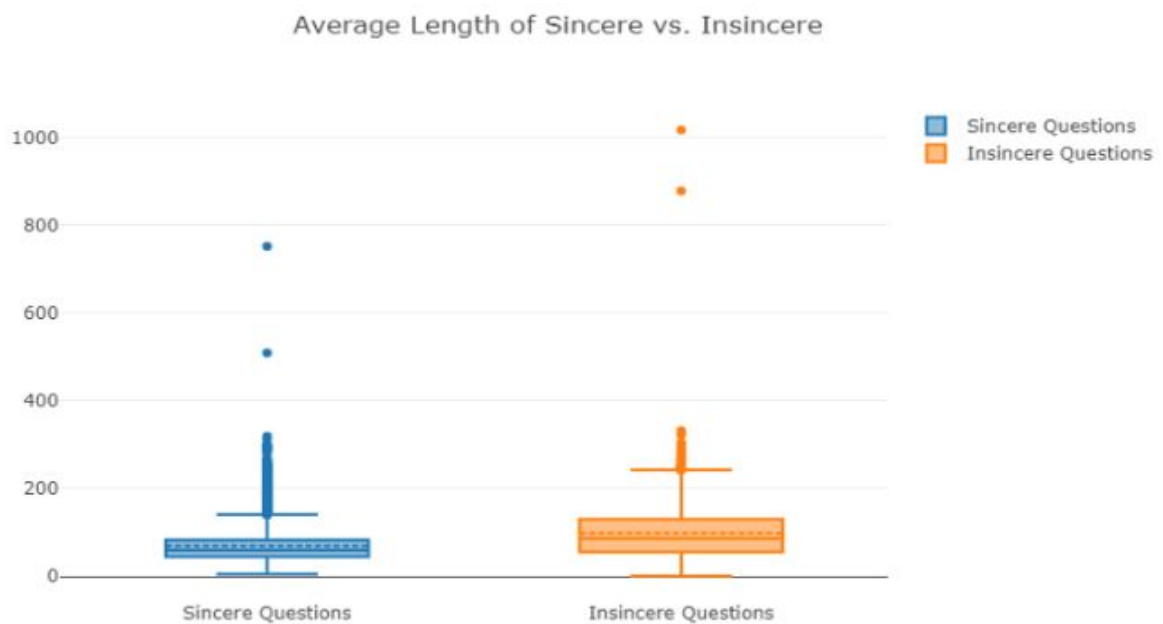
Exploratory Data Analysis:

The distribution of target:



We can see that this data set is highly imbalance with only 6.19 percent of insincere question and 93.81 percent of sincere question. Resampling and data augmentation maybe improve model performance. Moreover, evaluation metric F1-score will be work in this case because it considers both precision and recall of the test to compute the score.

The length of question:



Seems like length doesn't explain insincerity but we can see that the length of insincere questions are greater than sincere questions. Let's check some maximum and minimum length question:

Question: What is $\overbrace{\sum_{\vartheta=8}^{\infty} \overbrace{\frac{\sum_{\kappa=7}^{\infty} \overbrace{1x^0}^{\text{Read carefully.}} - 3x^{-1}} \div 1x^5 + \{\sqrt[3]{2x^{-3}}\}^{\{1x^0 + \vec{\vec{\{3x^{-3}\}^{\{1x^{-2}\}}\}}\}} \sum_{\dagger=9}^{\infty} \vec{\boxed{\boxed{3x^{-1}} + 3x^1 \times 1x^{-5}}\}} \div \sin(\boxed{\boxed{\vec{3x^{-5}}}} + \sqrt[4]{2x^{-4}} + \vec{2x^{-3}} \div \sin(\sqrt[5]{\int_{1x^5}^{2x^5} 2x^{-3} d\varrho)} \times \vec{\underbrace{2x^1}_{\text{Prove This.}}\}^{\{3x^4\}} \div \sqrt[5]{2x^{-3}} + \sum_{\theta=8}^{\infty} 1x^4\}} \times \{\boxed{\vec{\sum_{\nu=8}^{\infty} \sum_{4=6}^{\infty} \sum_{\xi=9}^{\infty} \boxed{3x^1} - \boxed{\sqrt[3]{\sqrt[3]{2x^{-2}}\}}\}}\}^{\{1x^3\}} - \cos(\{\tan(\sum_{0=6}^{\infty} 1x^{-3} \div 1x^1)\}^{\text{Molar Quantity.}})\}^{\{1x^3\}}^{\{1x^{-4}\}}\}^{\text{Expanded.}} [/math]?$

Length of question: 878

Target: 1

Question: What is $\frac{\int_{1x^5}^{3x^{-5}} \tan(\tan(\{\boxed{\int_{1x^0}^{1x^2} \sum_{\varpi=1}^{\infty} \int_{2x^{-3}}^{2x^2} \sum_{\alpha=7}^{\infty} \underbrace{\sqrt[2]{1x^5}}_{\text{Gauss's Law of Theoretical Probability.}} d\tau dx\}^{\{1x^0\}}) d\mu \{\int_{2x^{-3}}^{1x^5} \cos(\int_{2x^2}^{1x^{-3}}$

$$\frac{\sqrt{2}}{\overbrace{\underbrace{\frac{3x^3+3x^5}{\sqrt{3}2x^{-3}}}}_{\text{Gauss's Law of Theoretical Probability.}}} \times \overbrace{\tan(2x^0)}^{\text{Gauss's Law of Theoretical Probability.}} - \sum_{4=7}^{\infty} \boxed{3x^{-5}}^{\text{Inverse Function.}} \{ \boxed{\int_{2x^2}^{2x^4} 3x^1 dx} \div \sum_{6=6}^{\infty} \sqrt{3}2x^2 + \sqrt{4} \sin(2x^0+3x^0) \}^{2x^{-4}} + \boxed{\frac{\vec{\boxed{\sum_{\gamma=10}^{\infty} 1x^{-5}}}}{\frac{\sum_{\iota=2}^{\infty} 1x^{-5} - \frac{3x^{-1}}{1x^{-4}} \sin(\tan(3x^{-2}))}}}} \times \boxed{\sqrt{2} \{ \{ \sqrt{5}2x^5 \}^{2x^{-1}} \}^{2x^{-1}} \} \div \sum_{\chi=6}^{\infty} \int_{1x^4}^{2x^{-4}} 3x^2 d\vartheta + \{ 2x^{-3} \}^{2x^{-5}} \}^{3x^{-4}} \} d\mu d\iota} [/math]$$

Length of question: 1017

Target: 1

Question: "

Length of question: 1

Target: 1

Question: W

Length of question: 1

Target: 1

Question: "

Length of question: 1

Target: 1

Question: W

Length of question: 1

Target: 1

Question: Is

Length of question: 3

Target: 1

Question: If

Length of question: 3

Target: 1

Question: Do

Length of question: 3

Target: 1

Question: Is

Length of question: 3

Target: 1

Question: If
Length of question: 3
Target: 1

Question: Do
Length of question: 3
Target: 1

Question: Is
Length of question: 3
Target: 1

Question: If
Length of question: 3
Target: 1

Question: Do
Length of question: 3
Target: 1

- **There is significant noise in our data. Almost questions that have only one word are labeled insincere.**
- **Besides, some questions have 2 words isn't noise:**

Question: What's ejection?
Number of word in question: 2
Target: 0

Question: What's WhatsApp?
Number of word in question: 2
Target: 0


Question: What cyber?
Number of word in question: 2
Target: 0

Question: Google jigsaw?
Number of word in question: 2
Target: 0

Question: Sykes-Picot Agreement?
Number of word in question: 2
Target: 0

Check most used words in each class of questions:

Top used word in sincere question



Top used word in insincere question



Observations:

- Some of the top used words are common across both the classes likes 'think', 'india', ...
- The other top used words in sincere question after excluding the common ones are 'one', 'make', 'good', 'best', ...
- The other top used words in sincere question after excluding the common ones are 'one', 'make', 'good', 'best', ...

Text Pre-processing:

- Punctuation removal
- Stopword removal
- Tokenization

Model Selection

1, Try machine learning with Logistic Regression.

Feature extraction: Using TF-IDF.

Model: Logistic Regression CV.

Parameters:

- interception = True
- penalty = l2 norm
- solver = lbfgs
- max iteration = 20
- random state = 1

Result after training on test set:

F1 score = 0.5485900545785324

Accuracy score = 0.9544071202985932

2, Using deep learning with BiLSTM.

Embedding: Using glove 300d set.

Padding to the end of question with max length = 70.

Model Summary:

Layer (type)	Output Shape	Param #
input_20 (InputLayer)	[(None, 70)]	0
embedding_19 (Embedding)	(None, 70, 300)	300000000
bidirectional_45 (BidirectionalLSTM)	(None, 70, 256)	440320
bidirectional_46 (BidirectionalLSTM)	(None, 70, 128)	164864
bidirectional_47 (BidirectionalLSTM)	(None, 70, 64)	41472
dropout_37 (Dropout)	(None, 70, 64)	0
flatten_13 (Flatten)	(None, 4480)	0
dense_29 (Dense)	(None, 512)	2294272
dropout_38 (Dropout)	(None, 512)	0
dense_30 (Dense)	(None, 1)	513
Total params: 32,941,441		
Trainable params: 2,941,441		
Non-trainable params: 30,000,000		

- First layer is embedding layer using weights of glove set and output shape is (n_samples, max_length, embedding dim)
- Second, 3th and 4th layer is three Bidirectional LSTM layers with sequence output dimensional is: 256, 128, 64
- After three BiLSTM layer, we using one Drop out layer to decrease connected of previor layer with probability = 0.2.
- Then we Flatten output of Drop out layer to shape (n_sample, dimensional)
- Next, we using 1 fully connected layer with relu as activation and output dimensional: 512
- We using Dropout layer again with prob = 0.2.
- Last, we using one Fully connected layer with sigmoid as activation function with output dimensional: 1.

Training model:

- *Loss Function: Cross Entropy Loss*
- *Optimization: Adam*
- *Training Metric: Accuracy*
- *Validation set = 0.2, Training set = 0.8*
- *Number of epochs: 10*

- *Batch size = 512.*

Evaluation model:

- Evaluation metric: F1 score.
- We find the best threshold that give the max f1 score in range (0.1, 0.501) with step 0.01.

Result on training set:

```
Train on 835917 samples, validate on 208980 samples
835917/835917 [=====] - 147s 175us/sample - loss: 0.1330 - acc: 0.9497 - val_loss: 0.1152 - val_acc: 0.9551
Val F1 Score: 0.6218, Best Threshold: 0.3600
Train on 835917 samples, validate on 208980 samples
835917/835917 [=====] - 144s 173us/sample - loss: 0.1105 - acc: 0.9568 - val_loss: 0.1087 - val_acc: 0.9572
Val F1 Score: 0.6446, Best Threshold: 0.3600
Train on 835917 samples, validate on 208980 samples
835917/835917 [=====] - 144s 172us/sample - loss: 0.1034 - acc: 0.9592 - val_loss: 0.1064 - val_acc: 0.9569
Val F1 Score: 0.6522, Best Threshold: 0.4200
Train on 835917 samples, validate on 208980 samples
835917/835917 [=====] - 144s 172us/sample - loss: 0.0975 - acc: 0.9612 - val_loss: 0.1086 - val_acc: 0.9584
Val F1 Score: 0.6560, Best Threshold: 0.2800
Train on 835917 samples, validate on 208980 samples
835917/835917 [=====] - 144s 172us/sample - loss: 0.0910 - acc: 0.9636 - val_loss: 0.1073 - val_acc: 0.9584
Val F1 Score: 0.6583, Best Threshold: 0.2500
Train on 835917 samples, validate on 208980 samples
835917/835917 [=====] - 144s 172us/sample - loss: 0.0834 - acc: 0.9663 - val_loss: 0.1106 - val_acc: 0.9555
Val F1 Score: 0.6564, Best Threshold: 0.4400
```

There is significant overfitting after 6 epochs when loss on validation set at 6th epoch greater than loss at 5th epoch while loss on training set continue decrease and model early stop after 6 epochs.

Result on Testing set:

- F1-Score: 0.6623 with threshold 0.49
- Accuracy score: 0.9569 with threshold 0.5

FURTHER IMPROVEMENT:

- Using more time for Exploratory Data Analysis to understand more about data.
- Try resampling and augmentation to balane data.
- Try some other embedding set.
- Try some new model architecture.
- Target to archive F1 score more than 0.8 on test set.

Contribution of each member in group:

1, Nguyễn Bá Nghĩa:

- *Exploratory Data Analysis.*
- *Text Preprocessing.*

2, Nguyễn Thành Long:

- *Model selection.*
- *Embedding.*
- *Evaluate model and improvement.*